

1 Question 1

1.1

See code.

1.2

Device Bandwidth GB/sec			
	char	uint	uint2
Problem Size MB			
1.23515	24.3216	79.9968	93.1205
2.4703	48.0977	106.773	114.706
4.9406	54.0027	125.523	133.964
9.8812	57.7983	137.33	146.867
19.7624	60.5199	143.689	154.703
39.5248	61.4533	146.068	159.539
79.0496	61.8805	149.951	162.067
158.099	62.1447	151.075	163.309
316.198	62.2718	144.983	163.82

1.3

Considering the theoretical peak bandwidth is 240GB/s, but that for the char kernel it is estimated to be at most 25%, this gives us an expected bandwidth of about 60GB/s. We see roughly this bandwidth for the char kernel once the problem size exceeds 19.7624 MB. For the smaller problem sizes the execution time is likely being dominated by read and write duration and the capacity of the GPU is being underutilized.

The speed up from switching to uint happens because the number of reads and writes is reduced. For every 4 chars we only need to read one uint, the execution time for the threads when using chars and using uints is most likely very similar since we are only doing a single addition. This is why for the smaller problem sizes we see a roughly 4x speedup. For the larger problem sizes we cannot load all the text data into the registers on the GPU at once, so we start to run against the memory bandwidth limits.

For the uint2, at smaller sizes, we get a further speedup but still are dominated by the read and write times. However for larger problems since we have 8 times fewer read and writes, we approach the benchmark speed of the K80 on Google cloud.

2 Question 2

2.1

See code.

2.2

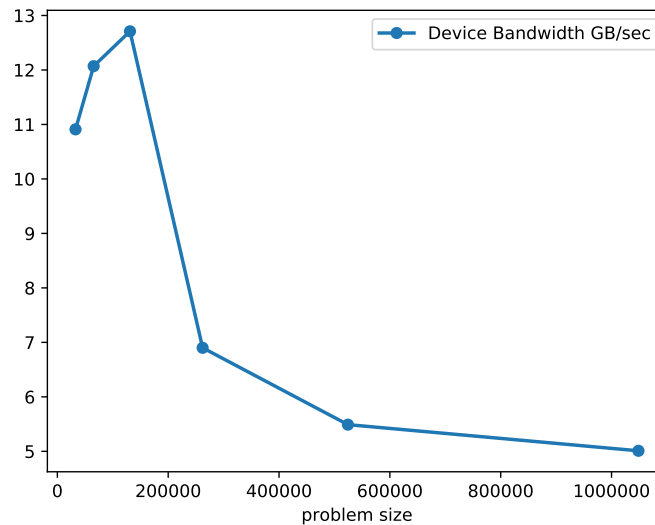
For every call to `device_graph_propagate` we write one float and read 1 unsigned integer and 2 floats for every edge of the current node.

Therefore the total number of bytes read is

$$total_bytes = node(f + edge(u + 2f))NUM_ITERATIONS \quad (1)$$

where f is the size of a float and u is the size of an unsigned integer in bytes.

2.3



2.4

For this problem, compared to problem 2, we are doing many more reads for each kernel call. The number of reads increases as the average edge size increases. This is reflected in the bandwidth measurements, where bandwidth decreases as average number of edges increases.

For example, with an average number of edges we are doing 31 reads and we get slightly more than $1/31$ of the theoretical peak performance.

	Device Bandwidth GB/sec					
	Number of nodes					
Avg. no. edges	32768	65536	131072	262144	524288	1048576
2	12.12	14.27	15.22	9.59	7.09	6.31
3	11.91	13.61	14.67	9.25	6.80	6.03
4	12.23	13.44	14.45	8.77	6.54	5.81
5	12.05	13.22	14.12	8.39	6.31	5.64
6	10.08	11.10	11.60	8.11	5.78	5.19
7	10.03	10.87	11.39	7.89	5.58	5.08
8	10.95	11.68	12.57	7.45	5.68	5.18
9	10.71	11.72	12.13	7.17	5.59	5.06
10	10.91	12.07	12.71	6.90	5.49	5.01
11	10.84	11.86	12.43	6.68	5.42	4.93
12	10.55	11.54	12.15	6.39	5.28	4.82
13	10.13	10.98	11.83	6.20	5.16	4.74
14	10.21	11.01	11.55	6.07	5.06	4.67
15	10.02	11.00	11.27	5.89	4.97	4.59
16	9.95	10.87	11.00	5.75	4.87	4.51
17	9.43	10.42	10.72	5.61	4.80	4.47
18	9.52	10.36	10.54	5.46	4.72	4.39
19	9.28	10.36	10.28	5.37	4.65	4.34